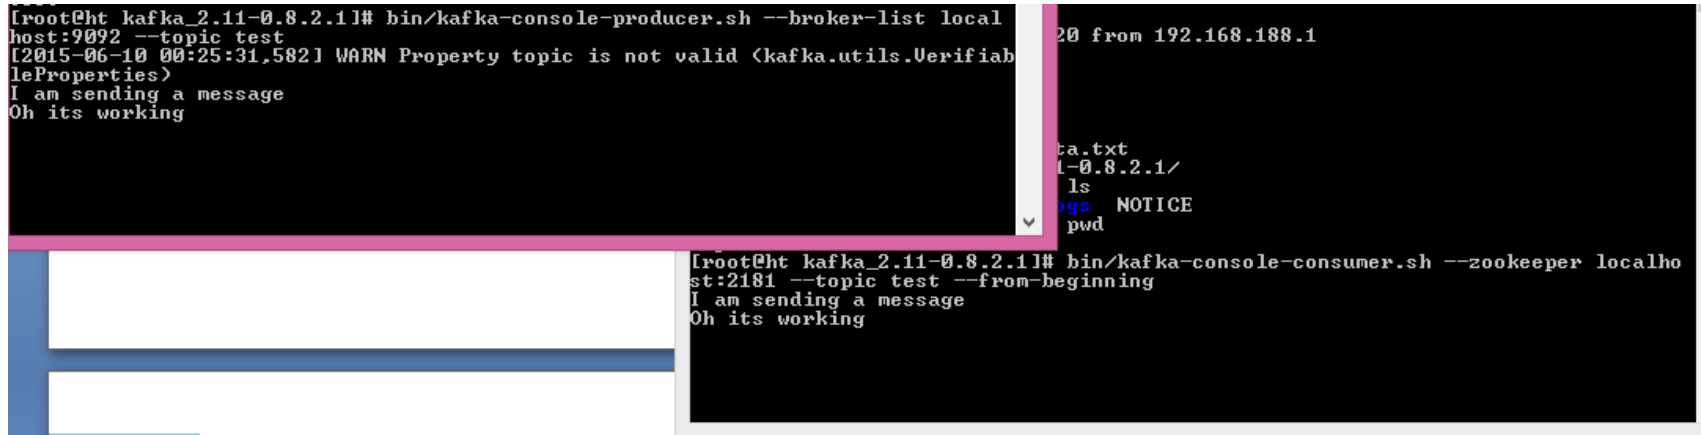


Try sending some information from the producer to consumer.



The image shows two terminal windows side-by-side. The left window shows the execution of the Kafka console producer, and the right window shows the execution of the Kafka console consumer. Both windows show successful execution of the commands.

```
[root@ht kafka_2.11-0.8.2.1]# bin/kafka-console-producer.sh --broker-list local
host:9092 --topic test
[2015-06-10 00:25:31,582] WARN Property topic is not valid (kafka.utils.Verifiab
leProperties)
I am sending a message
Oh its working
```

```
[root@ht kafka_2.11-0.8.2.1]# bin/kafka-console-consumer.sh --zookeeper localho
st:2181 --topic test --from-beginning
I am sending a message
Oh its working
```

You have successfully configure Kafka.

----- Kafka Installation ends Here -----

Miscellaneous:

Quit Scala CLI :- :q

Error: No suitable device found: no device found for connection 'System eth0'.

You could try removing/renaming `/etc/udev/rules.d/70-persistent-net.rules` and reboot. A new file will be created, sometimes that fixes this kind of problems.

You could also add a file `"/etc/sysconfig/network-scripts/ifcfg-eth1"`, can you bring up eth1 then?

Changing hostname:

- `/etc/hosts`
- `vi /etc/sysconfig/network`
- `sysctl kernel.hostname=slave`
- `bash`

Unable to start spark shell with the following error

```
at
org.apache.hadoop.hive.ql.session.SessionState.createSessionDirs(SessionState.java:554)
  at org.apache.hadoop.hive.ql.session.SessionState.start(SessionState.java:508)
  ... 85 more
Caused by: java.net.ConnectException: Connection refused
  at sun.nio.ch.SocketChannelImpl.checkConnect(Native Method)
  at sun.nio.ch.SocketChannelImpl.finishConnect(SocketChannelImpl.java:717)
  at org.apache.hadoop.net.SocketIOWithTimeout.connect(SocketIOWithTimeout.java:206)
```

```

at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:531)
at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:495)
at org.apache.hadoop.ipc.Client$Connection.setupConnection(Client.java:614)
at org.apache.hadoop.ipc.Client$Connection.setupIOstreams(Client.java:712)
at org.apache.hadoop.ipc.Client$Connection.access$2900(Client.java:375)
at org.apache.hadoop.ipc.Client.getConnection(Client.java:1528)
at org.apache.hadoop.ipc.Client.call(Client.java:1451)

```

Solution: unset HADOOP_CONF_DIR and unset HADOOP_HOME

<https://jaceklaskowski.gitbooks.io/mastering-apache-spark/>

issue: Cannot assign requested address

17/03/15 21:44:15 ERROR spark.SparkContext: Error initializing SparkContext.
 java.net.BindException: Cannot assign requested address: Service 'sparkDriver' failed after 16 retries (starting from 0)! Consider explicitly setting the appropriate port for the service 'sparkDriver' (for example spark.ui.port for SparkUI) to an available port or increasing spark.port.maxRetries.

```

    at sun.nio.ch.Net.bindo(Native Method)
    at sun.nio.ch.Net.bind(Net.java:433)
    at sun.nio.ch.Net.bind(Net.java:425)
    at sun.nio.ch.ServerSocketChannelImpl.bind(ServerSocketChannelImpl.java:
223)

```

Solution:

- include host to ip entry in /etc/hosts --> 192.168.188.178 master
- update in conf/spark-env.sh
 - SPARK_LOCAL_IP=127.0.0.1
 - SPARK_MASTER_HOST=192.168.188.178
- comment the following in conf/spark-defaults.conf
 - # spark.master spark://master:7077

Unable to start spark shell:

Caused by: java.io.IOException: Error accessing /opt/jdk/jre/lib/ext/_cld*.jar
 Such kind of error resolve by removing all the hidden file. .* (You can determine the file by running #ls -alt). List all the hidden dot file and remove it.

Yum repo config

```
# mkdir -p /redhatimg
# mount -o loop,ro /mnt/hgfs/MyExperiment/rhel-server-6.4-x86_64-dvd.iso /redhatimg
```

Create an entry in /etc/fstab so that the system always mounts the DVD image after a reboot.

```
/mnt/hgfs/MyExperiment/rhel-server-6.4-x86_64-dvd.iso /redhatimg iso9660 loop,ro 0 0
```

```
/etc/yum.repos.d/rhdiso.repo
```

```
[rhel6dvdiso]
name=RedHatOS DVD ISO
mediaid=1359576196.686790
baseurl=file:///redhating
enabled=1
gpgcheck=1
gpgkey=file:///etc/pki/rpm-gpg/RPM-GPG-KEY-redhat-release
```

Scala IDE

Download Scala Eclipse or install scala plugin from Eclipse Marketplace.

Eclipse IDE - [Eclipse IDE for Enterprise Java Developers](#)

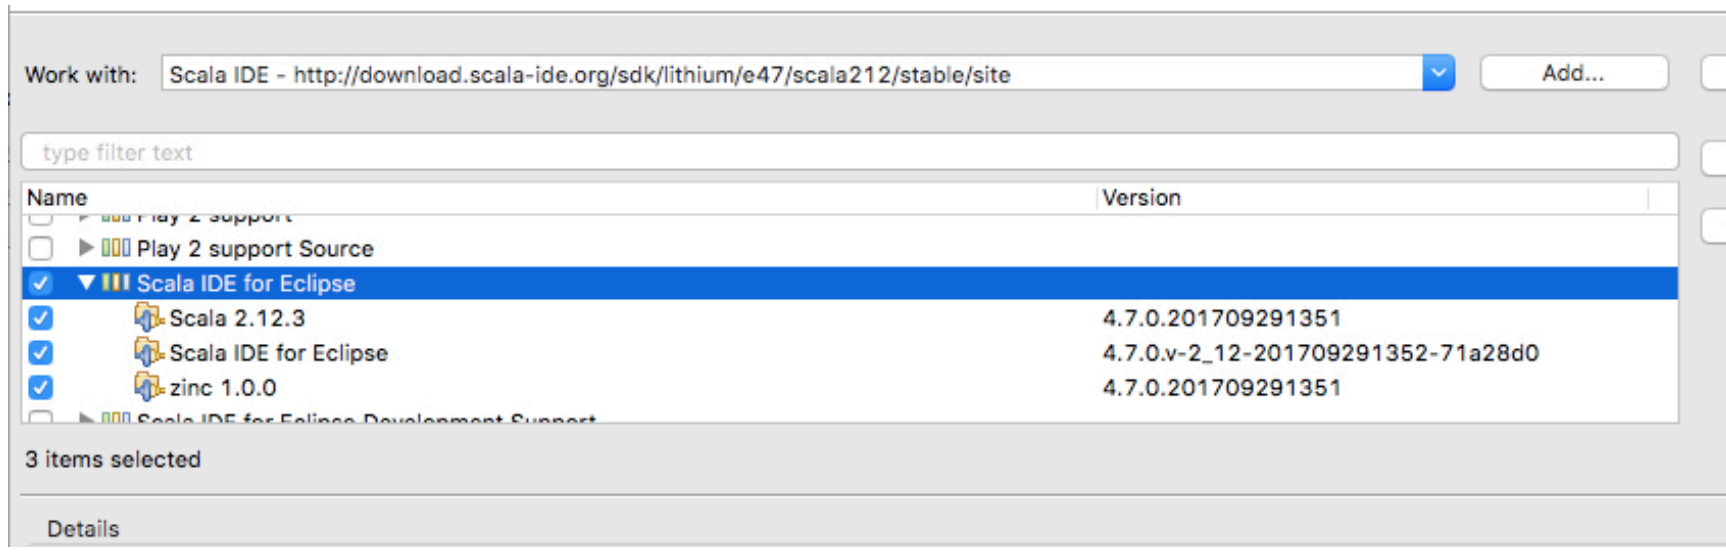
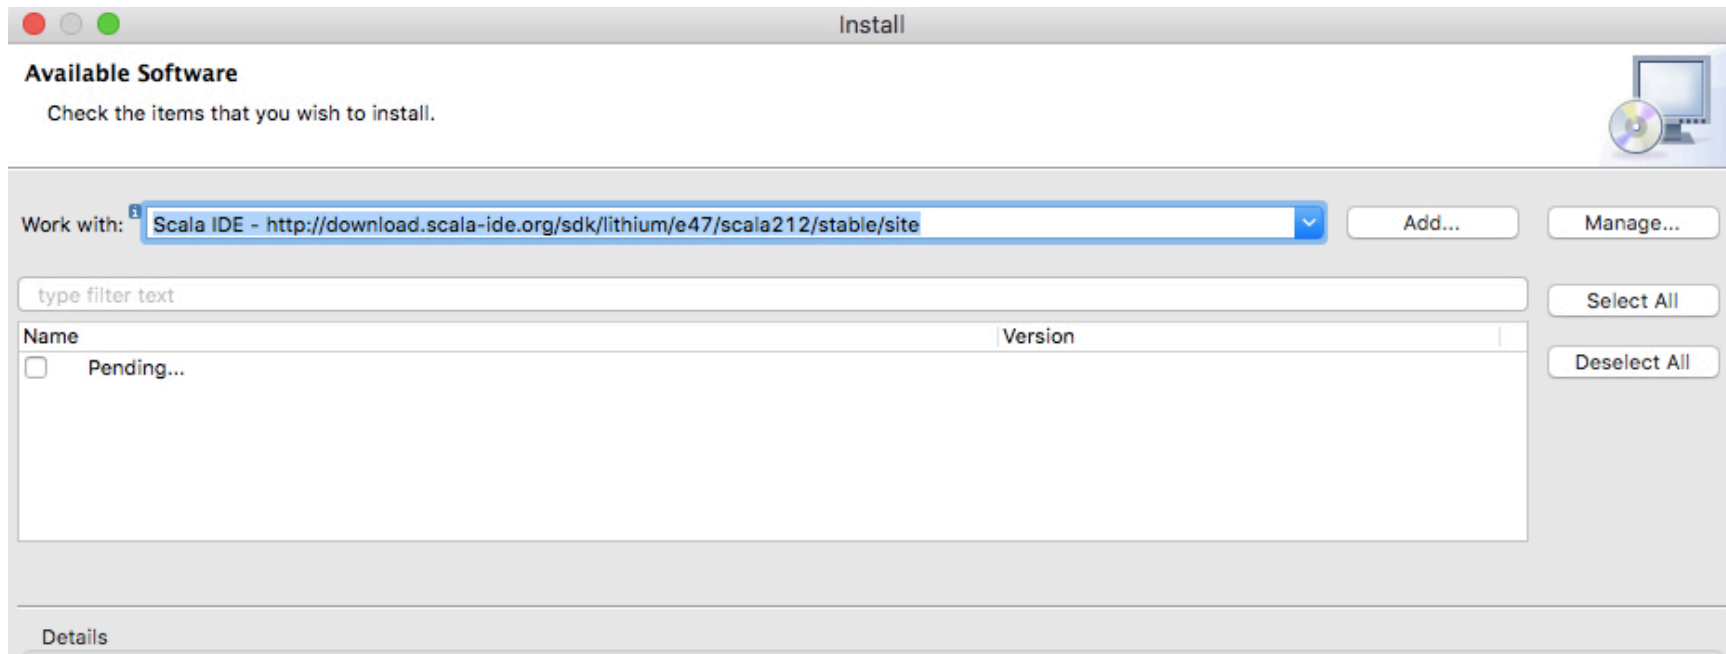
<https://www.eclipse.org/downloads/packages/>

Scala IDE for Eclipse is best installed (and updated) directly from within Eclipse.

<http://download.scala-ide.org/sdk/lithium/e47/scala212/stable/site>

This is done by using **Help** → **Install New Software...**, add the **Add...** button in the dialog.

Choose a name for the update site (Scala IDE is an obvious choice). Then read the next section to select which version you will install.



<http://scala-ide.org/download/current.html>

Or Download IDE

<http://scala-ide.org/download/sdk.html>

Download Latest Spark : Specify in the first Lab.

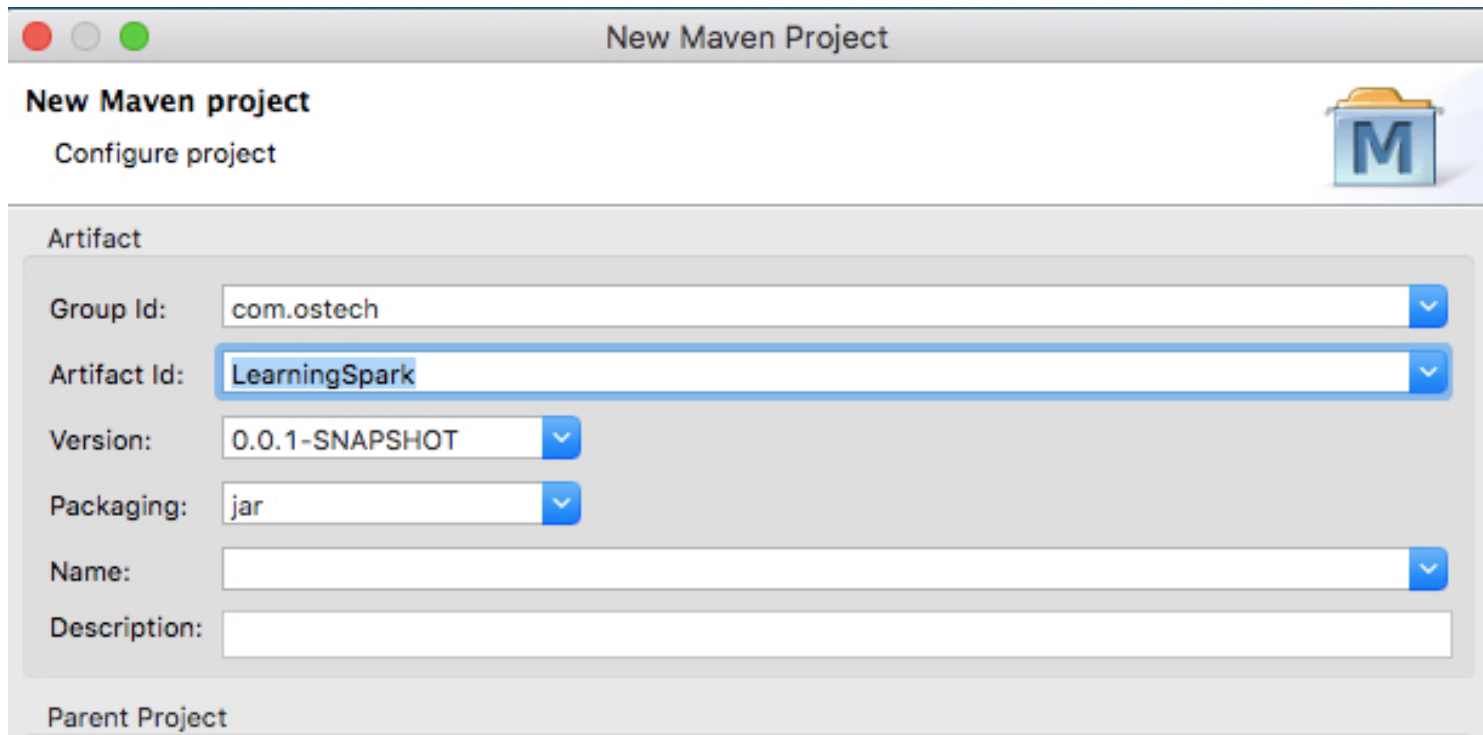
[\[https://spark.apache.org/downloads.html\]](https://spark.apache.org/downloads.html).

Add Spark Libraries

Go to `Java Build Path`, and add all the jars present under `scala-n.n.n-bin-hadoopN.N/jars/`.

Create a simple maven project:

Com.ostech:LearningSpark



Adding packages in Zeppelin

If using zeppelin, configure the following in the Interpreter.

spark.jars.packages

org.apache.spark:spark-streaming-kafka:0.10.2,12.3.0.1

Comma-separated list of Maven coordinates of jars to include on the driver and executor classpaths. The coordinates should be groupid:artifactid:version. If spark.jars.ivySettings is given artifacts will be resolved according to the configuration in the file, otherwise artifacts will be searched for in the local maven repo, then maven central and finally any additional remote repositories given by the command-line option --repositories.